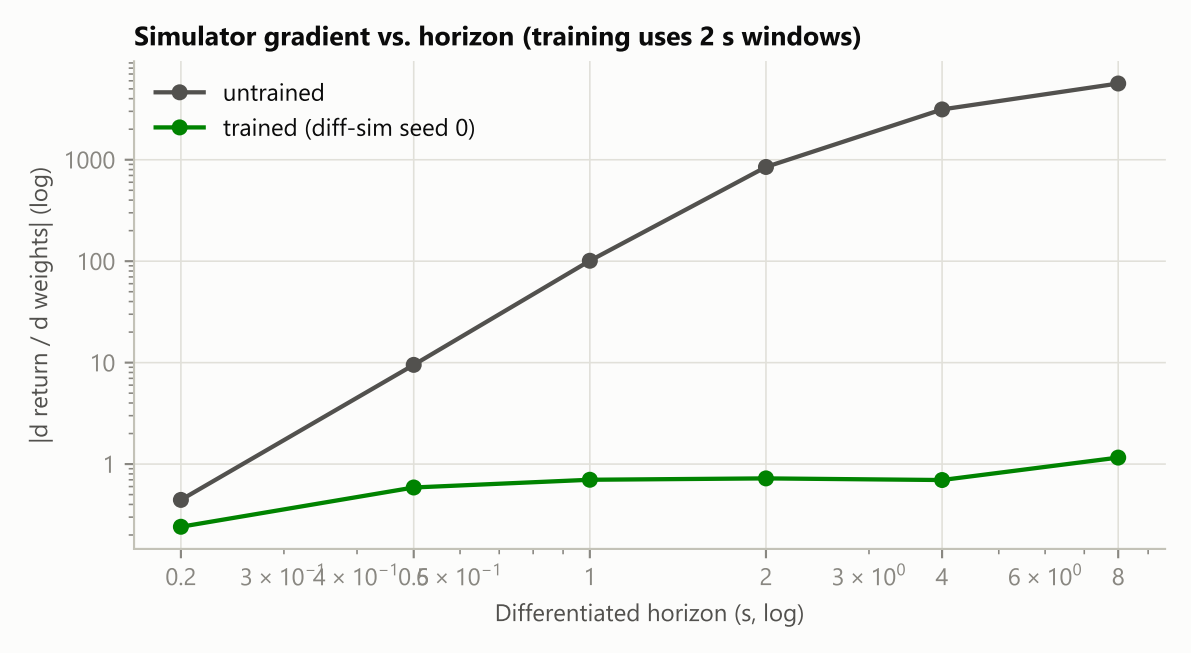
Simulator gradient vs. horizon (training uses 2 s windows)
untrained
trained (diff-sim seed 0)
|d return / d weights| (log)
1000
100
10
1
0.2
3 × 10⁻¹
4 × 10⁻¹
0.6 × 10⁻¹
1
2
3 × 10⁰
4
6 × 10⁰
8
Differentiated horizon (s, log)